



UNIVERSIDAD DE GRANADA / INSTITUTO DE ASTROFÍSICA DE ANDALUCÍA

BIGDATA: INTEGRATING NoSQL TECHNOLOGIES INTO VIRTUAL OBSERVATORY

MÁSTER EN MÉTODOS Y TÉCNICAS AVANZADAS EN FÍSICA
TRABAJO FIN DE MÁSTER PRESENTADO POR JOSÉ ANTONIO MAGRO CORTÉS

2013

Contents

1	Overview	1
2	Astronomy Big Data	3
2.1	Atacama Large Millimeter Array	4
2.1.1	ALMA Instrument	4
2.1.2	ALMA Science Archive	7
2.2	Square Kilometer Array	8
2.2.1	Figures and facts	9
3	The Virtual Observatory	12
3.1	ObsTAP	13
3.1.1	ObsCore Data Model	14
3.1.2	Table Access Protocol	14
3.2	Flexible Image Transport System	15
3.3	OpenCADC	17
3.3.1	Universal Worker Service	18
3.3.2	cadctap Library	20
4	Relational DBMS vs Document oriented DBMS	22
4.1	Relational Database Management System	22
4.1.1	Postgresql	22
4.2	NoSQL	25
4.2.1	MongoDB: a document oriented database	28
4.2.2	Advantages and uncertainties of using a NoSQL solution	31
4.3	A FITS alternative with MongoDB	34

5	Rewriting ASA with MongoDB	35
6	Successful experiences	36
7	Conclusions and future work	37
7.0.1	Conclusions	37
7.0.2	Future work	37
A	Getting and installing MongoDB	38
B	Configuring Eclipse and NetBeans for Java/Mongo development	40
	Bibliography	42

Chapter 1

Overview

The volume of data produced at some science centers presents a considerable processing challenge.

At CERN, for instance, almost 600 million times per second, particles collide within the Large Hadron Collider (LHC). Each collision generates particles that often decay in complex ways into even more particles. Electronic circuits record the passage of each particle through a detector as a series of electronic signals, and send the data to the CERN Data Centre for digital reconstruction. The digitized summary is recorded as a "collision event". Physicists must sift through the 15 petabytes or so of data produced annually to determine if the collisions have thrown up any interesting physics.

CERN does not have the computing or financial resources to crunch all of the data on site, so in 2002 it turned to grid computing to share the burden with computer centers around the world. The Worldwide LHC Computing Grid is a distributed computing infrastructure arranged in tiers which gives a community of over 8000 physicists near real-time access to LHC data..

So we could think that current deployed technologies are maybe obsolete and a new rethink should be made. If CERN (but not just CERN as we will discuss later) has changed its storage policy (leaving behind the classic client-server model), why

not change the way the data are accessed?

Almost a decade after E. Codd published his famous relational model paper, the relational database management systems (RDBMS) have been the must-be tools. Today, non-relational, cloud, or the so-called NoSQL databases are gaining mindshare as an alternative model for database management. NoSQL movement is growing rapidly. Often more characteristics apply such as schema-free, easy replication support, simple API, eventually consistent / BASE (not ACID), a huge amount of data (big data) and more.

In this document, we present an alternative, using one of the NoSQL databases free available, to offer a different way of challenging these problems, and always operating inside VO frame.

Chapter 2

Astronomy Big Data

Big data are high-volume, high-velocity, and/or high-variety information assets that require new forms of processing to enable enhanced decision making, insight discovery and process optimization. Examples of big data include information such as Web search results, electronic messages (e.g. SMS, email and instant messages), social media postings, pictures, videos and even system log data. However, it can also include cash transactions, check images, receipts and other transactional information depending on the source of the information.

Astronomical datasets are growing at an exponential rate: high performance computing applications in astronomy are enabling complex simulations with many billions of particles, while the forthcoming generation of telescopes will collect data at rates in excess of terabytes per day. This data deluge, both now and into the future, presents some critical challenges for the way astronomers derive new knowledge from their data.

In this chapter, we make an overview of some of the greatest telescopes, the way they acquire and store data and the problem they face.

2.1 Atacama Large Millimeter Array

2.1.1 ALMA Instrument

ALMA is a worldwide project; the synthesis of early visions of astronomers in its three partner communities: Europe, North America and Japan.

ALMA is a fusion of ideas, with its roots in three astronomical projects: the Millimeter Array (MMA) of the United States, the Large Southern Array (LSA) of Europe, and the Large Millimeter Array (LMA) of Japan.



Figure 2.1: ALMA dishes in Atacama Desert

Millimeter Array

The origins of the Millimeter Array (MMA) are found in the science of the NRAO 36-Foot Telescope (later known as the 12-Meter Telescope), soon followed by the 4.9m telescopes at the University of Texas and Aerospace Corporation, the 14m telescope at the Five Colleges Radio Astronomical Observatory, and the 7m telescope at AT&T Bell Labs. The science targets of the MMA included the same broad range of topics

seen at the VLA: (NSF) in July, 1990, called for an array of 40 antennas of 8-meter diameter, with four receiver bands covering the atmospheric windows from 30-350 GHz, configurable in four arrays of size 70-3000 m. The proposal discussed two possible sites for the MMA, both in the southwestern United States. Studies of the atmospheric transparency and phase stability at these sites led to similar studies on Mauna Kea, in Hawaii.

Concerns with the limited size of the area available to the MMA on Mauna Kea and with potential environmental problems prompted a search for potential sites in Chile. From April 1994, many high-elevation sites were visited. Finally, the site retained for the MMA was formally proposed in 1996: it was the Chajnantor plateau.

This effort culminated in the signing of the ALMA Agreement on February 25, 2003, between the North American and European parties. More than 14 government agencies in Chile were involved in the negotiations.

Assuming all three partners are able to meet their commitments, it was decided that the final project would be cost-shared 37.5% / 37.5% / 25% between North America, Europe, and Japan, respectively. The observing time, after a 10% share for Chile, would be shared accordingly.

ALMA is an instrument that consists of a giant array of 12-m antennas with baselines up to 16 km, and an additional compact array of 7-m and 12-m antennas to greatly enhance ALMA's ability to image extended targets, located on the Chajnantor plateau at 5000m altitude. Initially, it will observe at wavelengths in the range 3mm to 400m (84 to 720GHz). The antennas can be moved around, in order to form arrays with different distributions of baseline lengths. More extended arrays will give high spatial resolution, more compact arrays give better sensitivity for extended sources. In addition to the array of 12-m antennas, there is the Atacama Compact Array (ACA), consisting of twelve 7-m antennas and four 12-m antennas. This array will mostly stay in a fixed configuration and is used to image large scale structures that are not well sampled by the ALMA 12-m array.

The design of ALMA is driven by three key science goals:

- The ability to detect spectral line emission from CO in a normal galaxy like the Milky Way at a redshift of $z = 3$, in less than 24 hours
- The ability to image the gas kinematics in protostars and in protoplanetary disks around young Sun-like stars in the nearest molecular clouds ($150pc$)
- The ability to provide precise high dynamic range images at an angular resolution of $0.1arcsec$

ALMA delivers data cubes, of which the third axis is frequency. In this sense, the final data products are very much like that of an integral field unit with up to a million Spectral Pixels.

Interferometry

In order to obtain images, the raw visibility data need to be Fourier transformed. When ALMA is in full operations, this imaging step, as well as various calibration steps, will be done in the data reduction pipeline. Thus, fully calibrated data cubes will be delivered to the user.

Observing frequencies

The frequency range available to ALMA is divided into different receiver bands. Data can only be taken in one band at a time. These bands range from band 3, starting at $84GHz$, to band 10, ending at $\sim 950GHz$.

Field of view

The FWHM of the ALMA primary beam is $21''$ at $300GHz$, and scales linearly with wavelength (diffraction limit of a single 12-m antenna, as opposed to that of the whole array).

Spatial resolution

The spatial resolution of ALMA depends on the observing frequency and the maximum baseline of the array, following the $1.2x\lambda/D$ scaling. In the most compact

configurations ($\sim 150m$), resolutions range from $0.7''$ at $675GHz$ to $4.8''$ at $110GHz$.

Array configurations

The ALMA 12-m array will cycle from its most compact configuration, with maximum baselines of $\sim 150m$, to its most extended configuration, with maximum baselines of $\sim 16km$ (when completed), and back.

Spectral resolution

ALMA can deliver data cubes with up to 7680 frequency channels (spectral resolution elements). The width of these channels can range between $3.8kHz$ and $15.6MHz$, but the total bandwidth cannot exceed $8GHz$.

2.1.2 ALMA Science Archive

As stated in [8], the purpose of the ALMA archive is to provide services for:

- Persistent archiving and retrieval for observacional data.
- Observaction descrpitors.
- Datacubes produced by pipeline.
- Technical and environmental data.

And the key-point of the conceptual design of the ALMA Archive is to guarantee that three ALMA Regional Centres (ARCs) hold an identical copy of the archive at the Joint ALMA Observatory (JAO) in Santiago.

Relational denormalized database.

ALMA frontend archive is optimized for storage and preservation, not for data query and retrieval. ASA database is inspired from ObsCore, RADAMS and Hubble Legacy Archive plus Virtual Observatory Software:

- openCADC, which is used for database access and VO access protocol

- VOView, for Web components

2.2 Square Kilometer Array

Thousands of linked radio wave receptors will be located in Australia and in Southern Africa. Combining the signals from the antennas in each region will create a telescope with a collecting area equivalent to a dish with an area of about one square kilometre.

The Square Kilometer Array (SKA) will address fundamental unanswered questions about our Universe including how the first stars and galaxies formed after the Big Bang, how galaxies have evolved since then, the role of magnetism in the cosmos, the nature of gravity, and the search for life beyond Earth.

The SKA is a global science and engineering project led by the SKA Organisation, a not-for-profit company with its headquarters at Jodrell Bank Observatory, near Manchester, UK.

An array of dish receptors will extend into eight African countries from a central core region in the Karoo desert of South Africa. A further array of mid frequency aperture arrays will also be built in the Karoo. A smaller array of dish receptors and an array of low frequency aperture arrays will be located in the Murchison Radio-astronomy Observatory in Western Australia.

Construction is scheduled to start in 2016.



Figure 2.2: Artist’s impression of the SKA dishes. Credit: SKA Organisation/TD-P/DRAO/Swinburne Astronomy Productions

2.2.1 Figures and facts

The recent launch of the Murchison Widefield Array (MWA) - a radio telescope based in Western Australia’s Mid West - marked the start of an impressive flow of astronomical data that will be stored in the iVEC-managed Pawsey Center in Kensington.

According to Professor Andreas Wicenec, from The University of Western Australia node of the International Centre for Radio Astronomy Research (ICRAR), SKA has “now have more than 400 megabytes per second of MWA data streaming along the National Broadband Network from the desert 800 km away”.¹

The Murchison Widefield Array is the first Square Kilometre Array precursor to enter full operations, generating a vast torrent of information that needs to be stored for later retrieval by researchers.

¹<http://www.skatelescope.org/news/pawsey-centre/>

According to Professor Wicenec, “To store the Big Data the MWA produces, you’d need almost three 1 TB hard drives every two hours’. The technical challenge isn’t just in saving the observations but how you then distribute them to astronomers from the MWA team in far-flung places so they can start using it”.

There are currently two links between the data stores in Perth and MWA researchers at the Massachusetts Institute of Technology (MIT) in the United States and the Victoria University of Wellington in New Zealand. A future link to India another MWA partner will also be created.

The data are not obviously intended to be fully available for everybody at every-time: for instance, MIT researchers are interested in the early universe so filtering techniques to control what data is copied from the Pawsey Center archive to the MIT machines are used. By 2013, more than 150 TB of data had been transferred automatically to the MIT store, with a stream of up to 4 TB a day increasing that value.

MWA is producing so much information that it would be impossible to manually decide where to send what, which is where a sophisticated archiving system the open-source Next Generation Archive System (NGAS) comes in. NGAS was initially developed by Professor Wicenec at the European Southern Observatory (ESO) and later modified by the ICRAR team to meet the MWA data challenge.

In the NGAS one simply asks the system for what the data wanted and it either provides it from the local store or retrieves it from the full archive back in Perth.

About half of all MWA computing occurs on site in the Murchison, where signals from radio telescope antennas are combined and processed in a powerful system of computers called a correlator. What’s left to do in Perth is produce images, and manage storage and distribution by the archive system so MWA astronomers can analyze the collected data. Data travels down a dedicated 10 gigabit per second connection between the Murchison Radio-astronomy Observatory (MRO) and Geraldton, and the trip to Perth is completed on Australia’s new high-speed National

Broadband Network.

The MWA will store about 3 Petabytes at the Pawsey Center each year. Another section of the Pawsey Center will be a supercomputing facility that includes computing for Australia's other SKA precursor, CSIROs Australian Square Kilometre Array Pathfinder (ASKAP), and projects from geoscience and other computationally intensive fields.

To sum up, some relevant figures and facts about SKA:

- The data collected by the SKA in a single day would take nearly two million years to playback on an ipod.
- The SKA central computer will have the processing power of about one hundred million PCs.
- The SKA will use enough optical fibre to wrap twice around the Earth.
- The dishes of the SKA will produce 10 times the global internet traffic.
- The SKA will generate enough raw data to fill 15 million 64 GB iPods every day.
- The SKA will be so sensitive that it will be able to detect an airport radar on a planet 50 light years away.
- Analysts estimate the London Olympics was the most data-heavy yet - with some 60 Gbytes, the equivalent of 3,000 photographs, travelling across the network in the Olympic Park every second. This however is only equivalent to the data rate from about half a low frequency aperture array station in SKA phase one.

Chapter 3

The Virtual Observatory

The International Virtual Observatory Alliance (IVOA) was formed in 2002 with the aim to "facilitate the international coordination and collaboration necessary for the development and deployment of the tools, systems and organizational structures necessary to enable the international utilization of astronomical archives as an integrated and interoperating virtual observatory."¹ The IVOA now comprises programs from Argentina, Armenia, Australia, Brazil, Canada, China, France, Germany, Hungary, India, Italy, Japan, Russia, Spain, the United Kingdom, Ukraine, and the United States and inter-governmental organizations (ESA and ESO). Membership is open to other national and international programs according to the IVOA Guidelines for Participation.

The IVOA focuses on the development of standards and encourages their implementation for the benefit of the worldwide astronomical community. Working Groups are constituted with cross-program membership in those areas where key interoperability standards and technologies have to be defined and agreed upon. The Working Groups develop standards using a process modeled after the World Wide Web Consortium, in which Working Drafts progress to Proposed Recommendations and finally to Recommendations. Recommendations may ultimately be endorsed by the Virtual Observatory Working Group of Commission 5 (Astronomical Data) of the International Astronomical Union. The IVOA also has Interest Groups that dis-

¹<http://www.ivoa.net/>

cuss experiences using VO technologies and provide feedback to the Working Groups. Ad-hoc and permanent committees deal with specific scientific and procedural topics. Interaction with other scientific disciplines interested in data inter-operability is also pursued through dedicated Liaison Groups.

A chair is chosen from among the representatives and serves an eighteen-month term, normally preceded by an eighteen-month term as deputy chair. The Executive Committee meets 3-4 times a year and its members represent their respective programs and are expected to be in a position to commit resources targeted at the achievement of common goals. Decisions by the IVOA Executive Committee are reached by consensus. The IVOA holds two Interoperability Workshops each year. These meetings are opportunities for the IVOA Groups and Committees to have face-to-face discussions.

In this section we are not going to make a deep study of the Virtual Observatory techniques, technologies, protocols and interfaces, just those needed and selected to explain our proposal of including NoSQL into VO. Bearing in mind that the problem we are trying to address is data modelling, we will just make an overview of those aspects related with our work.

3.1 ObsTAP

A data model is a description of the objects represented by a computer system with their properties and relationships, a logical model detailing the decomposition of a complex dataset into simpler elements, like a collection of concepts and rules used in defining data model.

In 2011 IVOA proposed a new recommendation: Observation Data Model Core Components and its Implementation in the Table Access Protocol ². That document was intended to be a description of the interface which integrated the data modeling

²<http://www.ivoa.net/documents/ObsCore/20110502/PR-ObsCore-v1.0-20110502.pdf>

and data access aspects in a single service:

$$\boxed{\boxed{\text{ObsCore data model} + \text{Table Access Protocol} = \text{ObsTAP}}}$$

3.1.1 ObsCore Data Model

This is the first block for ObsTAP.

3.1.2 Table Access Protocol

This is second block for ObsTAP, and defines a Web service for accessing tables containing astronomical catalogues. TAP is the protocol which underlies in the process of posing a query against a data source (or several data sources). The result of a query is a table, usually a VOTable.³

Queries which use TAP protocol can be made through several clients, like:

- TOPCAT
- TAPHandle which operates fully within the Web browser
- The TAP shell, a command line interface to querying TAP servers, complete with metadata management and command line completion.
- The GAVO VOTable library, which allows embedding TAP queries in Python.

The types of queries implemented by TAP are:

- Data queries
- Metadata queries
- Virtual Observatory Support Interface (VOSI)

TAP includes support for multiple query languages, including queries specified using the Astronomical Data Query Language (ADQL [1]) and the Parameterised

³Support for VOTable output is mandatory, while other formats may be available.

Query Language (PQL, under development). Other query languages are also supported, and this mechanism allows developments outside the IVOA to be used without modifying the TAP specification.

Finally, it also includes support for both synchronous and asynchronous queries. Special support is provided for spatially indexed queries using the spatial extensions in ADQL.

```
SELECT centerAlpha, centerDelta, dateObs FROM lsw.plates
WHERE dateObs BETWEEN '1903-01-12' AND '1903-01-19';
```

ADQL Query for locating an unexpected object on the night sky between January 12th and 18th, 1903.

3.2 Flexible Image Transport System

Flexible Image Transport System (FITS) is an open standard defining a digital file format useful for storage, transmission and processing of scientific and other images. FITS is the most commonly used digital file format in astronomy. Unlike many image formats, FITS is designed specifically for scientific data and hence includes many provisions for describing photometric and spatial calibration information, together with image origin metadata.

The FITS format was first standardized in 1981; it has evolved gradually since then, and the most recent version (3.0) was standardized in 2008. FITS was designed with an eye towards long-term archival storage, and the maxim once FITS, always FITS represents the requirement that developments to the format must be backwards compatible.

A major feature of the FITS format is that image metadata is stored in a human-readable ASCII header, so that an interested user can examine the headers to investigate a file of unknown provenance. The information in the header is designed to

calculate the byte offset of some information in the subsequent data unit to support direct access to the data cells. Each FITS file consists of one or more headers containing ASCII card images (80 character fixed-length strings) that carry keyword/value pairs, interleaved between data blocks. The keyword/value pairs provide information such as size, origin, coordinates, binary data format, free-form comments, history of the data, and anything else the creator desires: while many keywords are reserved for FITS use, the standard allows arbitrary use of the rest of the name-space.

FITS is also often used to store non-image data, such as spectra, photon lists, data cubes, or even structured data such as multi-table databases. A FITS file may contain several extensions, and each of these may contain a data object. For example, it is possible to store x-ray and infrared exposures in the same file.

FITS support is available in a variety of programming languages that are used for scientific work, including C, C++, C#, Fortran, IGOR Pro, IDL, Java, LabVIEW, Mathematica, MatLab, Perl, PDL, Python, R, and Tcl. The FITS Support Office at NASA/GSFC maintains a list of libraries and platforms that currently support FITS.

Image processing programs such as ImageJ, GIMP, Photoshop, XnView and IrfanView can generally read simple FITS images, but frequently cannot interpret more complex tables and databases. Scientific teams frequently write their own code to interact with their FITS data, using the tools available in their language of choice. The FITS Liberator software is used by imaging scientists at the European Space Agency, the European Southern Observatory and NASA. The SAOImage DS9 Astronomical Data Visualization Application is available for many OSs, and handles images and headers.

FITS Data Format

A FITS file is comprised of parts called Header Data Units (HDU), being the first HDU called primary HDU or primary array. This array can contain a 1-999 dimensional array. A typical primary array could contain a 1D spectrum, 2D image or 3D data cube. Any number of HDU can follow the main array, and are called FITS

extensions. Currently, three different extensions can be defined:

- Image extension, a 0-9999 dimensional array of pixels, which begins with XTENSION = 'IMAGE'
- ASCII table extension which stores tabular data in ASCII formats. They begin with XTENSION = 'TABLE'
- Binary table extension stores tabular data in binary representation. Headers start with XTENSION = 'BINTABLE'

Besides, there are additional type of HDU called random groups, but only used for radio interferometry.




Figure 3.1: Viewing FITS header in Aladin

3.3 OpenCADC

This research used the facilities of the Canadian Astronomy Data Centre operated by the National Research Council of Canada with the support of the Canadian Space

Agency.

OpenCADC (Canadian Astronomy Data Center) is a Virtual Observatory, used in ALMA Science Archive (for this reason is included in this chapter) tool which comprises several projects ⁴:

3.3.1 Universal Worker Service

The Universal Worker Service (UWS) pattern defines how to build asynchronous (the client does not wait for each request to be fulfilled; if the client disconnects from the service then the activity is not aborted), stateful (the service remembers results of a previous activity), job-oriented (the rules for setting and arranging the parameters for a job is called Job Description Language- JDL) services.

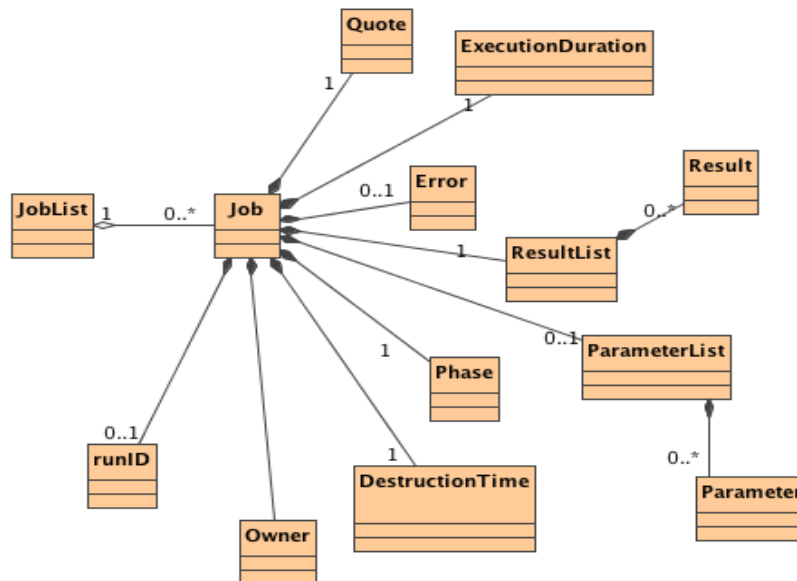


Figure 3.2: UWS Objects in class diagram

When the execution starts, the job is seen as a state machine which can adopt several states. The phases are the following:

⁴We do not enumerate all of them here, for a full list, go to <https://code.google.com/p/opencadc/source/browse>

- **PENDING**: the job is accepted but not committed for execution by the client.
- **QUEUED**: the job is committed for execution by the client but the service has not yet assigned it to a processor.
- **EXECUTING**: the job has been assigned to a processor and the results can be produced at any time during this phase.
- **COMPLETED**: execution is over and the user can collect the results.
- **ERROR**: the job failed to complete.
- **ABORTED**: the job has been stopped by the user or by the system.
- **UNKNOWN**: the job is in an unknown state.
- **HELD**: the job is pending execution.
- **SUSPENDED**: The job has been suspended by the system during execution. This might be because of temporary lack of resource. Similar to aborted but in this state the UWS will automatically resume the when resource are available again.

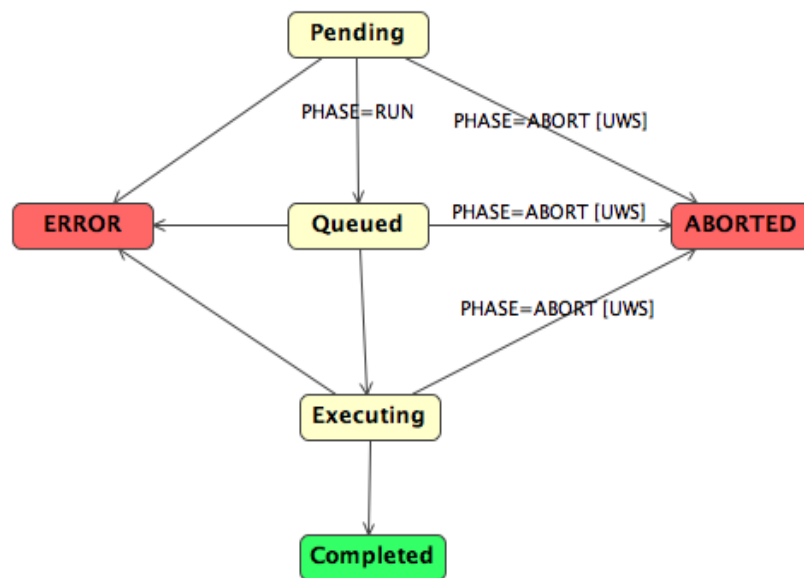


Figure 3.3: Universal Worker Server Job states

Now, we make a subtle description of the structure of the code relating UWS in OpenCADC (the cadcUWS Library), the section of the code provides Job class and plugin architecture, servlet with UWS async and sync behaviour.

JobManager

Responsible for job control.

JobPersistence

In charge of storing and retrieving Job state.

JobExecutor

It executes every job in separated threads.

JobRunner

The code that actually executes the job.

3.3.2 cadcTAP Library

The cadcTAP library is responsible of async and sync queries, where QueryRunner implements JobRunner. It also contains TAP_SCHEMA DDL statements and is used by query parser to validate table and column usage.

TapQuery Interface

It has a separate implementation for each LANG (e.g. $LANG = ADQL$) specified and processess the query to local SQL.

SqlQuery

When code states $LANG = SQL$, it implements TapQuery and fully navigates it (*FROM*, *WHERE* and *HAVING* clauses).

AdqlQuery

Same as before when $LANG = ADQL$.

Plugins

- UploadManager
- TableWriter
- FileStore

QueryRunner

It implements JobRunner and sets Job state, find DataSource and uses TapSchema, UploadManager, TapQuery and TableWriter.

Chapter 4

Relational DBMS vs Document oriented DBMS

Typically, modern relational databases have shown little efficiency in certain applications using intensive data, like indexing of a large number of documents, sites rendering with high traffic, and streaming sites. Typical RDBMS implementations are tuned either for small but frequent reads and writes or a large set of transactions that have few write accesses. On the other hand NoSQL can serve load lots of reads and writes.

We will start making a short overview of the relational model using Postgresql RDMS. Then we move to the NoSQL alternative, through Mongo database.

4.1 Relational Database Management System

ToDo

4.1.1 Postgresql

PostgreSQL is a powerful, open source object-relational database system. It has more than 15 years of active development and a proven architecture that has earned it a

strong reputation for reliability, data integrity, and correctness. It runs on all major operating systems, including Linux, UNIX (AIX, BSD, HP-UX, SGI IRIX, Mac OS X, Solaris, Tru64), and Windows. It is fully ACID compliant, has full support for foreign keys, joins, views, triggers, and stored procedures (in multiple languages). It includes most SQL:2008 data types, including INTEGER, NUMERIC, BOOLEAN, CHAR, VARCHAR, DATE, INTERVAL, and TIMESTAMP. It also supports storage of binary large objects, including pictures, sounds, or video. It has native programming interfaces for C/C++, Java, .Net, Perl, Python, Ruby, Tcl, ODBC, among others, and exceptional documentation.

An enterprise class database, PostgreSQL boasts sophisticated features such as Multi-Version Concurrency Control (MVCC), point in time recovery, tablespaces, asynchronous replication, nested transactions (savepoints), online/hot backups, a sophisticated query planner/optimizer, and write ahead logging for fault tolerance. It supports international character sets, multibyte character encodings, Unicode, and it is locale-aware for sorting, case-sensitivity, and formatting. It is highly scalable both in the sheer quantity of data it can manage and in the number of concurrent users it can accommodate. There are active PostgreSQL systems in production environments that manage in excess of 4 terabytes of data. Some general PostgreSQL limits are included in the table below.

PostgreSQL runs stored procedures in more than a dozen programming languages, including Java, Perl, Python, Ruby, Tcl, C/C++, and its own PL/pgSQL, which is similar to Oracle's PL/SQL. Included with its standard function library are hundreds of built-in functions that range from basic math and string operations to cryptography and Oracle compatibility. Triggers and stored procedures can be written in C and loaded into the database as a library, allowing great flexibility in extending its capabilities. Similarly, PostgreSQL includes a framework that allows developers to define and create their own custom data types along with supporting functions and operators that define their behavior. As a result, a host of advanced data types have been created that range from geometric and spatial primitives to network addresses to even ISBN/ISSN (International Standard Book Number/Inter-

national Standard Serial Number) data types, all of which can be optionally added to the system.

Just as there are many procedure languages supported by PostgreSQL, there are also many library interfaces as well, allowing various languages both compiled and interpreted to interface with PostgreSQL. There are interfaces for Java (JDBC), ODBC, Perl, Python, Ruby, C, C++, PHP, Lisp, Scheme, and Qt just to name a few.

Best of all, PostgreSQL's source code is available under a liberal open source license: the PostgreSQL License. This license gives you the freedom to use, modify and distribute PostgreSQL in any form you like, open or closed source. Any modifications, enhancements, or changes you make are yours to do with as you please. As such, PostgreSQL is not only a powerful database system capable of running the enterprise, it is a development platform upon which to develop in-house, web, or commercial software products that require a capable RDBMS.

Limit	Value
Maximum Database Size	Unlimited
Maximum Table Size	32 TB
Maximum Row Size	1.6 TB
Maximum Field Size	1 GB
Maximum Rows per Table	Unlimited
Maximum Columns per Table	250 - 1600 depending on column types
Maximum Indexes per Table	Unlimited

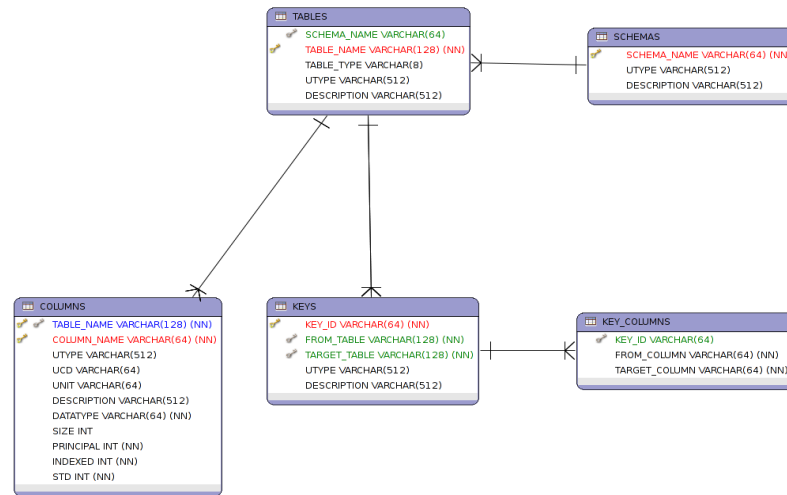


Figure 4.1: Entity-Relationship Diagram for VO Tap Schema

4.2 NoSQL

NoSQL implementations used in the real world include 3TB Digg green markers (indicated to highlight the stories voted by others in the social network), the 6 TB of "ENSEMBLE" European Commission database used in comparing models and air quality, and the 50 TB of Facebook inbox search.

NoSQL architectures often provide limited consistency, such as events or transactional consistency restricted to only data items. Some systems, however, provide all guarantees offered by ACID systems by adding an intermediate layer. There are two systems that have been deployed and provide for storage of snapshot isolation column: Google Percolator (based on BigTable system) and Hbase transactional system developed by the University of Waterloo. These systems use similar concepts in order to achieve distributed multiple rows ACID transactions with snapshot isolation guarantees for the underlying storage system in that column, with no extra overhead in data management, no system deployment middleware or any maintenance introduced by middleware layer.

Quite NoSQL systems employ a distributed architecture, maintaining data redundantly on multiple servers, often using distributed hash table. Thus, the system may actually escale adding more servers, and thus a server failure may be tolerated.

There are different NoSQL DBs for different projects:

- Document oriented
 - CouchDB
 - MongoDB
 - RavenDB
 - BaseX
 - djondb
 - eXist
 - SimpleDB
 - IBM Lotus Domino
 - Terrastore
- Graph oriented
 - Neo4j
 - DEX
 - AllegroGraph
 - OrientDB
 - InfiniteGraph
 - Sones GraphDB
 - InfoGrid
 - HyperGraphDB

- Key-value oriented
 - Cassandra
 - BigTable
 - Dynamo (Amazon)
 - MongoDB
 - Project Voldemort (LinkedIn)
 - Riak
 - Redis
- Multivalued
 - OpenQM
 - Extensible storage engine
- Object Oriented
 - Zope Object Database
 - db4o
 - GemStone S
 - Objectivity/DB
- Tabular
 - HBase
 - BigTable
 - LevelDB (BigTable open version)
 - Hypertable

They run on clusters of inexpensive machines.

4.2.1 MongoDB: a document oriented database

MongoDB (from "humongous") is an open source document-oriented database system developed and supported by 10gen. It is part of the NoSQL family of database systems. Instead of storing data in tables as is done in a "classical" relational database, MongoDB stores structured data as JSON-like documents with dynamic schemas (MongoDB calls the format BSON), making the integration of data in certain types of applications easier and faster.

10gen began Development of MongoDB in October 2007 and was not created to be just another database that tries to do everything for everyone. Instead, MongoDB was created to work with documents rather than rows, was extremely fast, massively scalable, and easy to use. In order to accomplish this, some features were excluded, namely support for transactions.

A document database is more like a collection of documents. Each entry is a document, and each one can have its own structure. If you want to add a field to an entry, you can do so without affecting any other entry.

Main features

- **Ad hoc queries**

MongoDB supports search by field, range queries, regular expression searches. Queries can return specific fields of documents and also include user-defined JavaScript functions.

- **Indexing**

Any field in a MongoDB document can be indexed (indices in MongoDB are conceptually similar to those in RDBMSes). Secondary indices are also available.

- **Replication**

MongoDB supports master-slave replication. A master can perform reads and writes. A slave copies data from the master and can only be used for reads or

backup (not writes). The slaves have the ability to select a new master if the current one goes down.

- **Load balancing**

MongoDB scales horizontally using sharding.[9] The developer chooses a shard key, which determines how the data in a collection will be distributed. The data is split into ranges (based on the shard key) and distributed across multiple shards. (A shard is a master with one or more slaves.) MongoDB can run over multiple servers, balancing the load and/or duplicating data to keep the system up and running in case of hardware failure. Automatic configuration is easy to deploy and new machines can be added to a running database.

- **File storage**

MongoDB could be used as a file system, taking advantage of load balancing and data replication features over multiple machines for storing files. This function, called GridFS,[10] is included with MongoDB drivers and available with no difficulty for development languages (see "Language Support" for a list of supported languages). MongoDB exposes functions for file manipulation and content to developers. GridFS is used, for example, in plugins for NGINX and lighttpd. In a multi-machine MongoDB system, files can be distributed and copied multiple times between machines transparently, thus effectively creating a load balanced and fault tolerant system.

- **Aggregation**

MapReduce can be used for batch processing of data and aggregation operations. The aggregation framework enables users to obtain the kind of results for which the SQL GROUP BY clause is used.

- **Server-side JavaScript execution**

JavaScript can be used in queries, aggregation functions (such as MapReduce), are sent directly to the database to be executed.

- **Capped collections**

MongoDB supports fixed-size collections called capped collections. This type

of collection maintains insertion order and, once the specified size has been reached, behaves like a circular queue.

Once we have seen the main features of MongoDB, we can move on to the language itself.

The basics

We must know four concepts to dig into MongoDB's world:

- Database: this concept is much like the RDBM counterpart.
- Collection: we can see a collection and a table as the same thing.
- Document: its equivalent in RDBM is the row, and a document is made up of fields.
- Field: is a lot like a column.

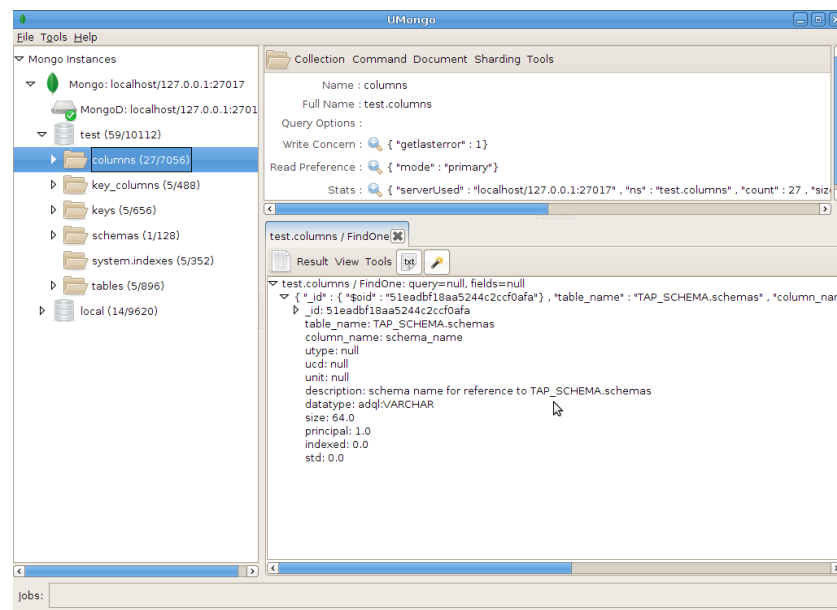


Figure 4.2: Tree View for Tap Schema in MongoVue

4.2.2 Advantages and uncertainties of using a NoSQL solution

- **Elastic scaling**

For years, database administrators have relied on scale up – buying bigger servers as database load increases – rather than scale out – distributing the database across multiple hosts as load increases. However, as transaction rates and availability requirements increase, and as databases move into the cloud or onto virtualized environments, the economic advantages of scaling out on commodity hardware become irresistible.

RDBMS might not scale out easily on commodity clusters, but the new breed of NoSQL databases are designed to expand transparently to take advantage of new nodes, and they're usually designed with low-cost commodity hardware in mind.

- **Big data**

Just as transaction rates have grown out of recognition over the last decade, the volumes of data that are being stored also have increased massively. O'Reilly has cleverly called this the industrial revolution of data. RDBMS capacity has been growing to match these increases, but as with transaction rates, the constraints of data volumes that can be practically managed by a single RDBMS are becoming intolerable for some enterprises. Today, the volumes of big data that can be handled by NoSQL systems, such as Hadoop, outstrip what can be handled by the biggest RDBMS.

- **No need for DBAs**

Despite the many manageability improvements claimed by RDBMS vendors over the years, high-end RDBMS systems can be maintained only with the assistance of expensive, highly trained DBAs. DBAs are intimately involved in the design, installation, and ongoing tuning of high-end RDBMS systems.

NoSQL databases are generally designed from the ground up to require less management: automatic repair, data distribution, and simpler data models lead to lower administration and tuning requirements – in theory. In practice, its

likely that rumors of the DBAs death have been slightly exaggerated. Someone will always be accountable for the performance and availability of any mission-critical data store.

- **Economics**

NoSQL databases typically use clusters of cheap commodity servers to manage the exploding data, while RDBMS tends to rely on expensive proprietary servers and storage systems. The result is that the cost per gigabyte or transaction/second for NoSQL can be many times less than the cost for RDBMS, allowing you to store and process more data at a much lower price point.

- **Flexible data models**

Change management is a big headache for large production RDBMS. Even minor changes to the data model of an RDBMS have to be carefully managed and may necessitate downtime or reduced service levels.

NoSQL databases have far more relaxed or even nonexistent data model restrictions. NoSQL Key Value stores and document databases allow the application to store virtually any structure it wants in a data element. Even the more rigidly defined BigTable-based NoSQL databases (Cassandra, HBase) typically allow new columns to be created without too much fuss.

The result is that application changes and database schema changes do not have to be managed as one complicated change unit. In theory, this will allow applications to iterate faster, though, clearly, there can be undesirable side effects if the application fails to manage data integrity.

NoSQL systems have generated a lot of enthusiasm but there are still a lot of questions about its future:

- **Maturity**

RDBMS systems have been around for a long time. NoSQL advocates will argue that their advancing age is a sign of their obsolescence, but for most CIOs, the maturity of the RDBMS is reassuring. For the most part, RDBMS systems are stable and richly functional. In comparison, most NoSQL alternatives are in

pre-production versions with many key features yet to be implemented. Living on the technological leading edge is an exciting prospect for many developers, but enterprises should approach it with extreme caution.

- **Support**

Enterprises want the reassurance that if a key system fails, they will be able to get timely and competent support. All RDBMS vendors go to great lengths to provide a high level of enterprise support.

In contrast, most NoSQL systems are open source projects, and although there are usually one or more firms offering support for each NoSQL database, these companies often are small start-ups without the global reach, support resources, or credibility of an Oracle, Microsoft, or IBM.

- **Analytics and business intelligence**

NoSQL databases have evolved to meet the scaling demands of modern Web 2.0 applications. Consequently, most of their feature set is oriented toward the demands of these applications. However, data in an application has value to the business that goes beyond the insert-read-update-delete cycle of a typical Web application. Businesses mine information in corporate databases to improve their efficiency and competitiveness, and business intelligence (BI) is a key IT issue for all medium to large companies.

NoSQL databases offer few facilities for ad-hoc query and analysis. Even a simple query requires significant programming expertise, and commonly used BI tools do not provide connectivity to NoSQL.

Some relief is provided by the emergence of solutions such as HIVE or PIG, which can provide easier access to data held in Hadoop clusters and perhaps eventually, other NoSQL databases. Quest Software has developed a product Toad for Cloud Databases that can provide ad-hoc query capabilities to a variety of NoSQL databases.

- **Administration**

The design goals for NoSQL may be to provide a zero-admin solution, but the

current reality falls well short of that goal. NoSQL today requires a lot of skill to install and a lot of effort to maintain.

- **Expertise**

There are literally millions of developers throughout the world, and in every business segment, who are familiar with RDBMS concepts and programming. In contrast, almost every NoSQL developer is in a learning mode. This situation will address naturally over time, but for now, it is by far easier to find experienced RDBMS programmers or administrators than NoSQL experts.

4.3 A FITS alternative with MongoDB

Chapter 5

Rewriting ASA with MongoDB

Esto es una prueba

Chapter 6

Successful experiences

Chapter 7

Conclusions and future work

7.0.1 Conclusions

- La ingente cantidad de datos recopilada en los proyectos de vanguardia se está convirtiendo en un problema.
- Los sistemas, protocolos y modelos tradicionales no son válidos en todos los casos.
- Cualquier modificación que se haga no deberá implementarse al margen del observatorio virtual.
- Las bases de datos NoSQL, y en concreto las orientadas a documentos (pero no exclusivamente estas), pueden simplificar el análisis y el diseño de un sistema, pero también aumentar el rendimiento del tratamiento de datos, por las razones expuestas en este trabajo.

7.0.2 Future work

- Utilizar métricas formales para la rescritura del código.
- Utilizar benchmarks para obtener datos precisos sobre la mejora del rendimiento.
- Decidir si MongoDB es un sistema gestor de base de datos adecuado.
- Decidir qué lenguaje usar para atacar a la base de datos seleccionada.

Appendix A

Getting and installing MongoDB

- Go to <http://www.mongodb.org/downloads> and select our OS version. Download it and decompress.
- Supposing we are in a Linux box and that we are using the latest release (in June 2013) execute the following command:

```
./mongod --rest --dbpath /opt/mongodb-linux-i686-2.4.5/bin/data/db/
```

- The server is now listening and waiting for connections, by default, in port 27017. As we have used the *-rest* modifier, we can point our browser to <http://localhost:28017> for http diagnostic access.

mongod debian

[List all commands](#) | [Replica set status](#)

Commands: [buildInfo](#) [cursorInfo](#) [features](#) [hostInfo](#) [isMaster](#) [listDatabases](#) [repSetGetStatus](#) [serverStatus](#) [top](#)

db version v2.4.5
 git hash: a2ddc68ba7c9ee17bfe69ed840383ec3506602b
 sys info: Linux bs-linux32.10gen.cc 2.6.21.7-2.fc8xen #1 SMP Fri Feb 15 12:39:36 EST 2008 i686 BOOST_LIB_VERSION=1_49
 uptime: 11 seconds

overview (only reported if can acquire read lock quickly)

time to get readlock: 0ms
 # databases: 1
 # Cursors: 0
 replication:
 master: 0
 slave: 0

clients


Client	OpId	Locking	Waiting	SecsRunning	Op	Namespace	Query	client	msg	progress
clientcursorsmon	5		{ waitingForLock: false }		0			:27017		
snapshothread	2		{ waitingForLock: false }		0			:27017		
webserv	7		{ waitingForLock: false }		0			:27017		
DataFileSync	0		{ waitingForLock: false }		0			:27017		
initandlisten	6		{ waitingForLock: false }	2002	local.startup_log		0.0.0.0:0			
TTLMonitor	3		{ waitingForLock: false }		0			:27017		

dbtop (occurrences/percent of elapsed)

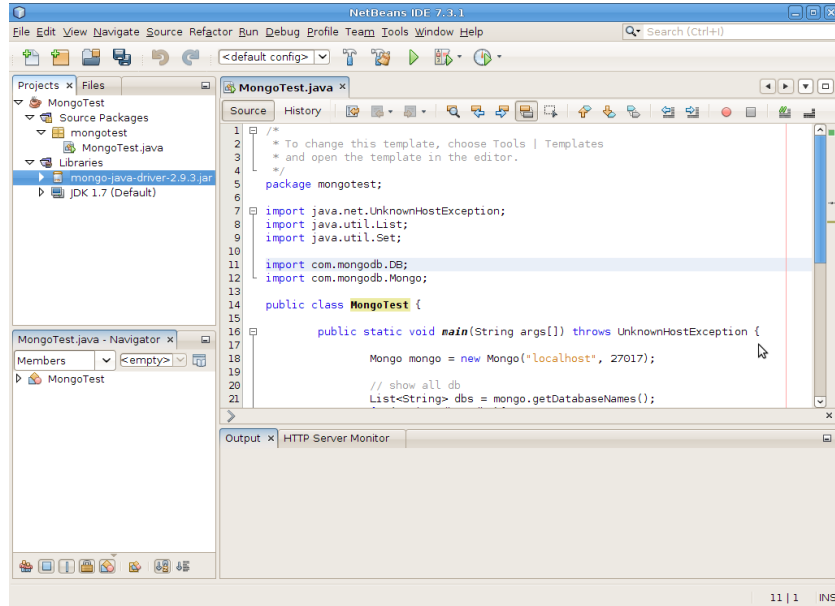
Appendix B

Configuring Eclipse and NetBeans for Java/Mongo development

- Go to <http://central.maven.org/maven2/org/mongodb/mongo-java-driver/> and choose the right version.
- In Eclipse, just start a new project **File - New - Java Project**
- Right-click in **Properties** and then **Java Build Path - Libraries - Add External JARs**



- Import MongoDB methods, classed and interfaces as needed.
- In NetBeans, start a new project **File - New Project - Java Application**
- Right-click in **Libraries** and then **Add JAR/Folder**



- Import MongoDB methods, classed and interfaces as needed.

Bibliography

- [1] Inc 10gen. The mongodb 2.4 manual. <http://docs.mongodb.org/manual/>, 2011-2013. [Online; accessed 01-December-2012].
- [2] Kristina Chodorow. *MongoDB: The Definitive Guide*. O'Reilly, 2010.
- [3] Kristina Chodorow. *50 Tips and Tricks for MongoDB Developers*. O'Reilly, 2011.
- [4] Juan de Dios Santander Vela. Bases de datos multimedia para radioastronomía: Radams y dss-63. <http://www.iaa.es/~jdsant/DEA-BBDDMultimediaParaRadioastronomia-RADAMS%2CDSS63.pdf>, 2006. [Online; accessed 01-May-2013].
- [5] Juan de Dios Santander Vela et al. The alma science archive: Design. <ftp://ftp.eso.org/projects/adass/posters/P147.pdf>, 2011. [Online; accessed 27-April-2013].
- [6] Juan de Dios Santander Vela et al. The alma science archive: Implementation. <ftp://ftp.eso.org/projects/adass/posters/P133.pdf>, 2011. [Online; accessed 07-February-2013].
- [7] Abraham Silberschatz et al. *Database System Concepts Sixth Edition*. McGraw-Hill, 2010.
- [8] Sandra Etoka. A first look at the alma science archive. <http://www.eso.org/sci/facilities/alma/meetings/gar-sep07/pdf/Etoka.pdf>, 2012. [Online; accessed 21-March-2013].

-
- [9] Centre for Astrophysics and Supercomputing. Scientific computing and visualisation. <http://astronomy.swin.edu.au/scivis/>, 2013. [Online; accessed 30-January-2013].
- [10] Centre for Astrophysics and Supercomputing. Accelerating astrophysics research with gpu supercomputing. <http://www.sgi.com/pdfs/4401.pdf>, 2013. [Online; accessed 30-January-2013].
- [11] Nature Vol 455 — Issue no. 7209. Big data: science in the petabyte era. <http://www.nature.com/nature/journal/v455/n7209/edsumm/e080904-01.html>, 2008. [Online; accessed 14-February-2011].
- [12] IBM Research. Square kilometer array: Ultimate big data challenge. <http://asmarterplanet.com/blog/2013/03/the-square-kilometer-array-the-world%E2%80%99s-ultimate-big-data-challenge.html>, 2013. [Online; accessed 04-July-2012].
- [13] John Richer. Alma, synthesis imaging and the astro-grid. <http://thames.cs.rhul.ac.uk/~fionn/astro-grid-papers/john-richer.pdf>, 2001. [Online; accessed 22-November-2012].
- [14] skatelescope.org. The age of astronomy big data is already here. <http://www.skatelescope.org/news/pawsey-centre/>, 2013. [Online; accessed 09-July-2012].
- [15] S. Tansley T. Hey and editors K. Toll. The fourth paradigm. data-intensive scientific discovery. http://research.microsoft.com/en-us/collaboration/fourthparadigm/4th_paradigm_book_complete_lr.pdf, 2009. [Online; accessed 31-March-2013].